

# Tagger: Practical PFC Deadlock Prevention in Data Center Networks


**Shuihai Hu\*(HKUST)**, Yibo Zhu, Peng Cheng, Chuanxiong Guo\*  
(Toutiao), Kun Tan\*(Huawei), Jitendra Padhye, Kai Chen (HKUST)  
Microsoft

CoNEXT 2017, Incheon, South Korea

# RDMA is Being Widely Deployed


## RDMA: Remote Direct Memory Access

- ❖ High throughput, low latency with low CPU overhead
- ❖ Microsoft, Google, etc. are deploying RDMA



# Priority Flow Control (PFC)


---



PAUSE upstream switch when PFC threshold reached

❖ Avoid packet drop due to buffer overflow


# A Simple Illustration of PFC Deadlock




**Due to Cyclic Buffer Dependency (CBD)  $A \rightarrow B \rightarrow C \rightarrow A$**   
**Not just a theoretical problem, we have seen it in our datacenters too!**

# CBD in the Clos Network

---




# CBD in the Clos Network




consider two flows initially follow shortest UP-DOWN paths

# CBD in the Clos Network




due to link failures, both flows are locally rerouted to non-shortest paths

# CBD in the Clos Network



these two **DOWN-UP bounced** flows create CBD



buffer dependency graph

**CBD: L2->S1->L3->S2->L2**



# Real in Production Data Centers?


---

Packet reroute measurements in more than 20 data centers:

**~100,000 DOWN-UP reroutes!**

# Handling Deadlock is Important

- **#1:** transient problem → **PERMANENT** deadlock
  - ❖ Transient loops due to link failures
  - ❖ Packet flooding
  - ❖ ...
- **#2:** small deadlock can cause large deadlock



# Three Key Challenges

---

What are the challenges in designing a practical deadlock prevention solution?

- No change to existing routing protocols or hardware
- Link failures & routing errors are unavoidable at scale
- Switches support at most 8 limited lossless priorities  
(and typically only two can be used)

# The Existing Deadlock Prevention Solutions

---

- #1: deadlock-free routing protocols
  - ❖ not supported by commodity switches (**fail challenge #1**)
  - ❖ not work with link failures or routing errors (**fail challenge #2**)
- #2: buffer management schemes
  - ❖ require a lot of lossless priorities (**fail challenge #3**)


Our answer: **Tagger**



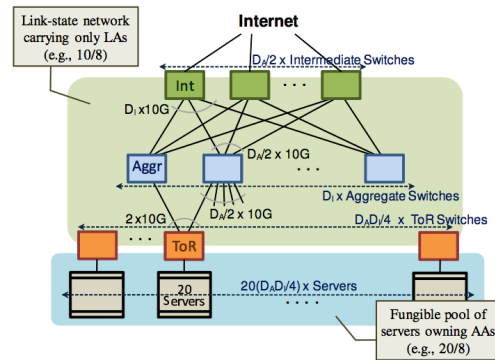
TAGGER DESIGN

# Important Observation


Fat-tree [Sigcomm'08]




VL2 [Sigcomm'09]



BCube [Sigcomm'09]



HyperX [SC'09]



**desired path set:** all shortest paths

**desired path set:** dimension-order paths


**Takeaway:** In a data center, we can ask operator to supply a set of **expected lossless paths (ELP)**!

# Basic Idea of Tagger

---

1. Ask operators to provide:
  - ❖ topology & expected lossless paths (ELP)
2. Packets carrying tags when in the network
3. Pre-install match-action rules at switches for tag manipulation and packet queueing
  - ❖ packets travel over ELP: lossless queues & CBD never forms
  - ❖ packets deviate ELP: lossy queue, thus PFC not triggered

# Illustrating Tagger for Clos Topology

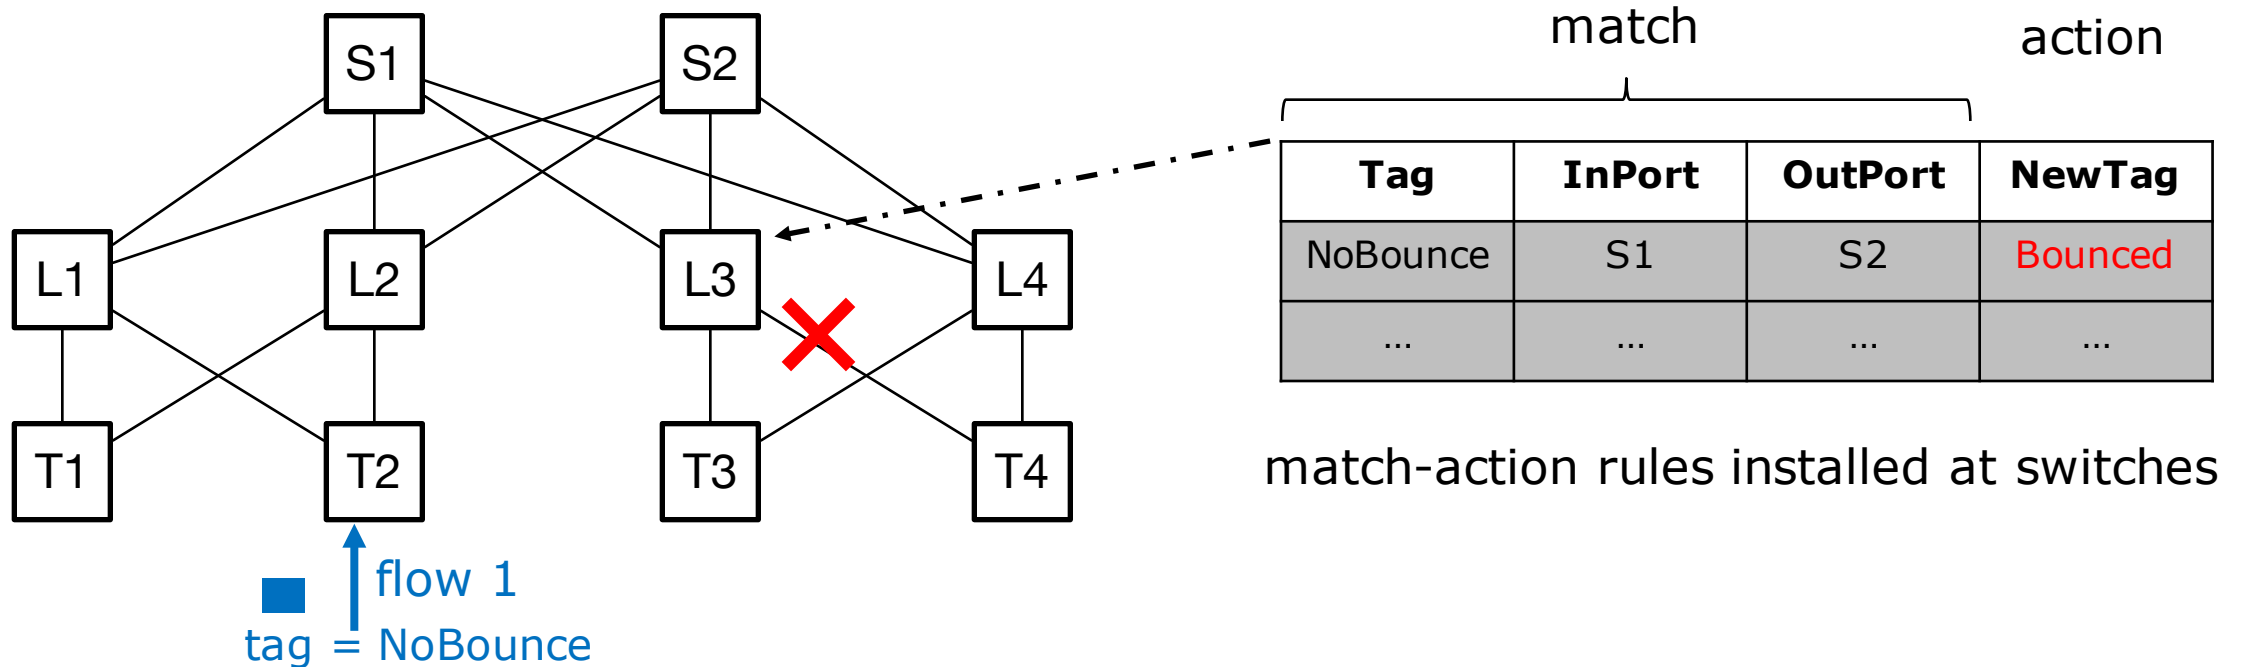


Root cause of CBD:  
packets deviate UP-DOWN routing!

**ELP** = all shortest paths (CBD-free)

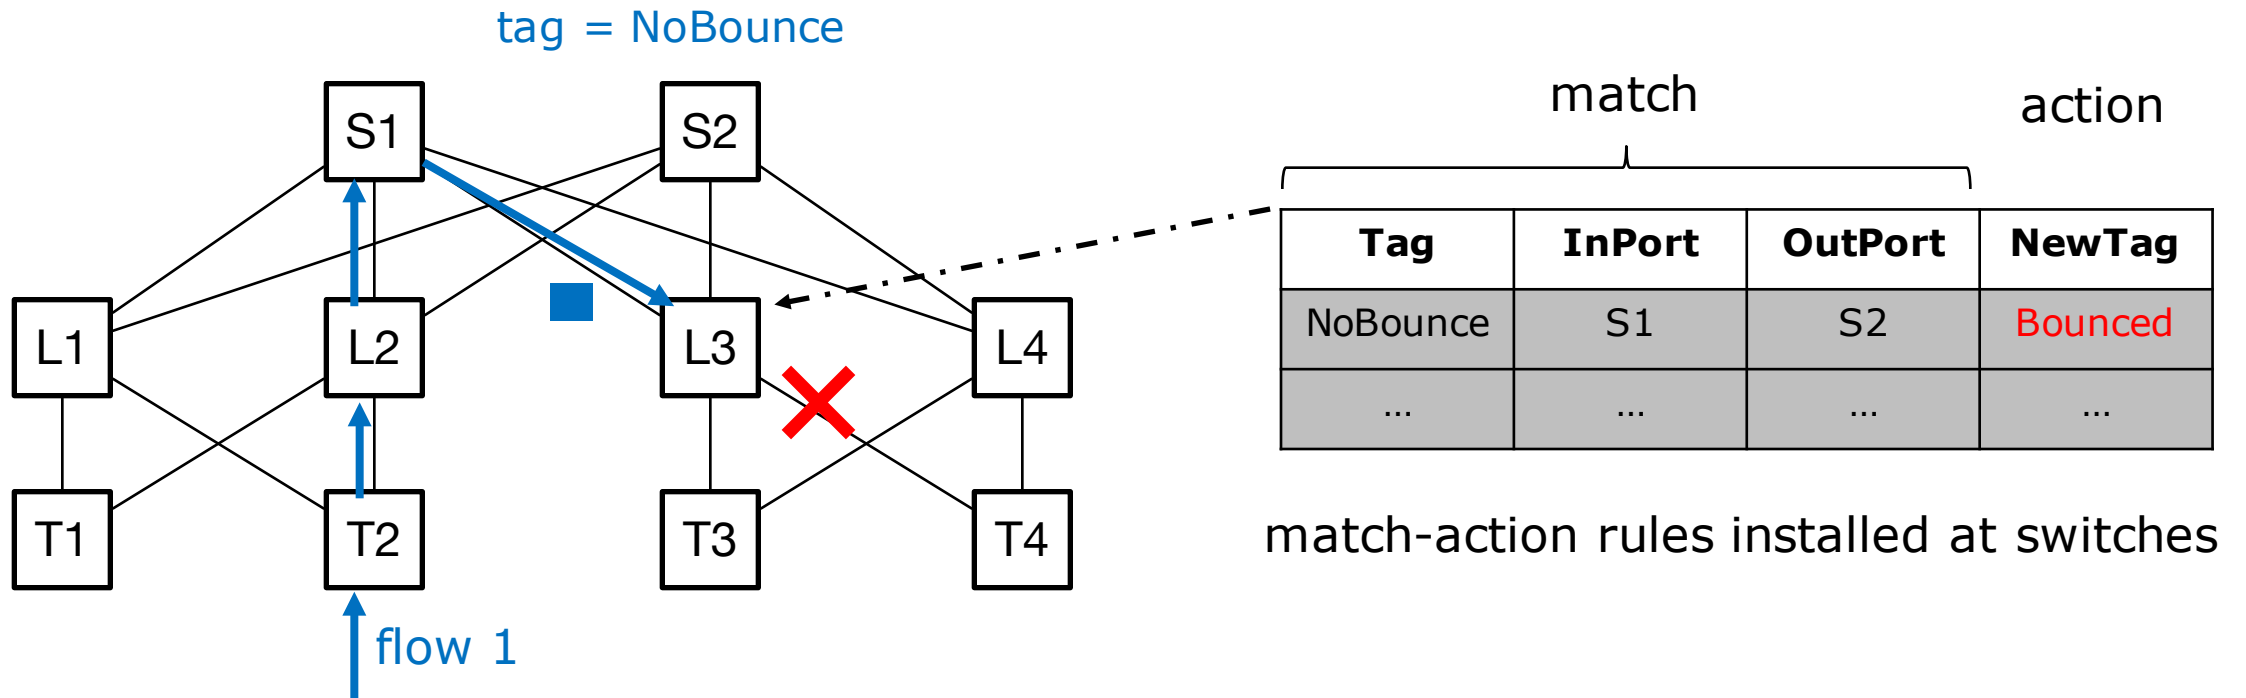


# Illustrating Tagger for Clos Topology



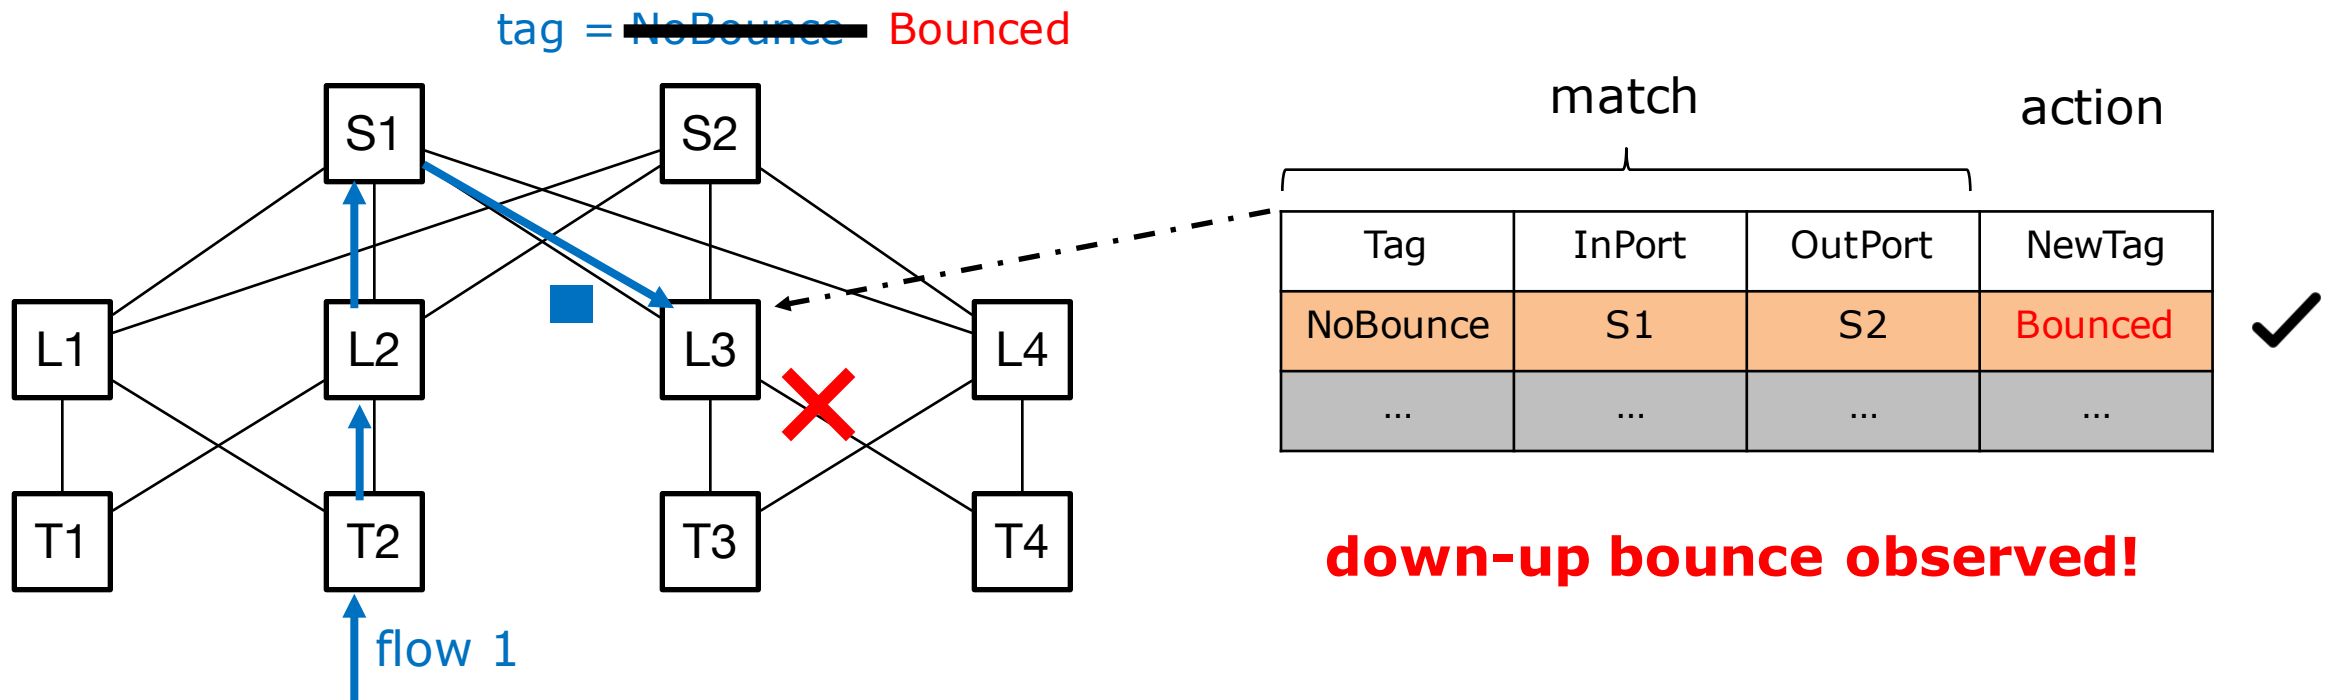
- Under Tagger, packets carry tags when travelling in the network
- Initially, tag value = NoBounce
- At switches, Tagger pre-install match-action rules for tag manipulation

# Illustrating Tagger for Clos Topology




Packet received by switch L3

# Illustrating Tagger for Clos Topology





rewrite tag once DOWN-UP bounce detected

# Illustrating Tagger for Clos Topology



- S2 knows it is a **bounced packet** that deviates ELP → **placed in the lossy queue**
- No PFC PAUSE sent from S2 to L3 → buffer dependency from L3 to S2 removed

# Illustrating Tagger for Clos Topology




buffer dependency graph

- Tagger will do the same for packets of flow 2
- 2 buffer dependency edges are removed → CBD is eliminated

~~CBD: L2 > S1 > L3 > S2 > L2~~

# What If ELP Has CBD?

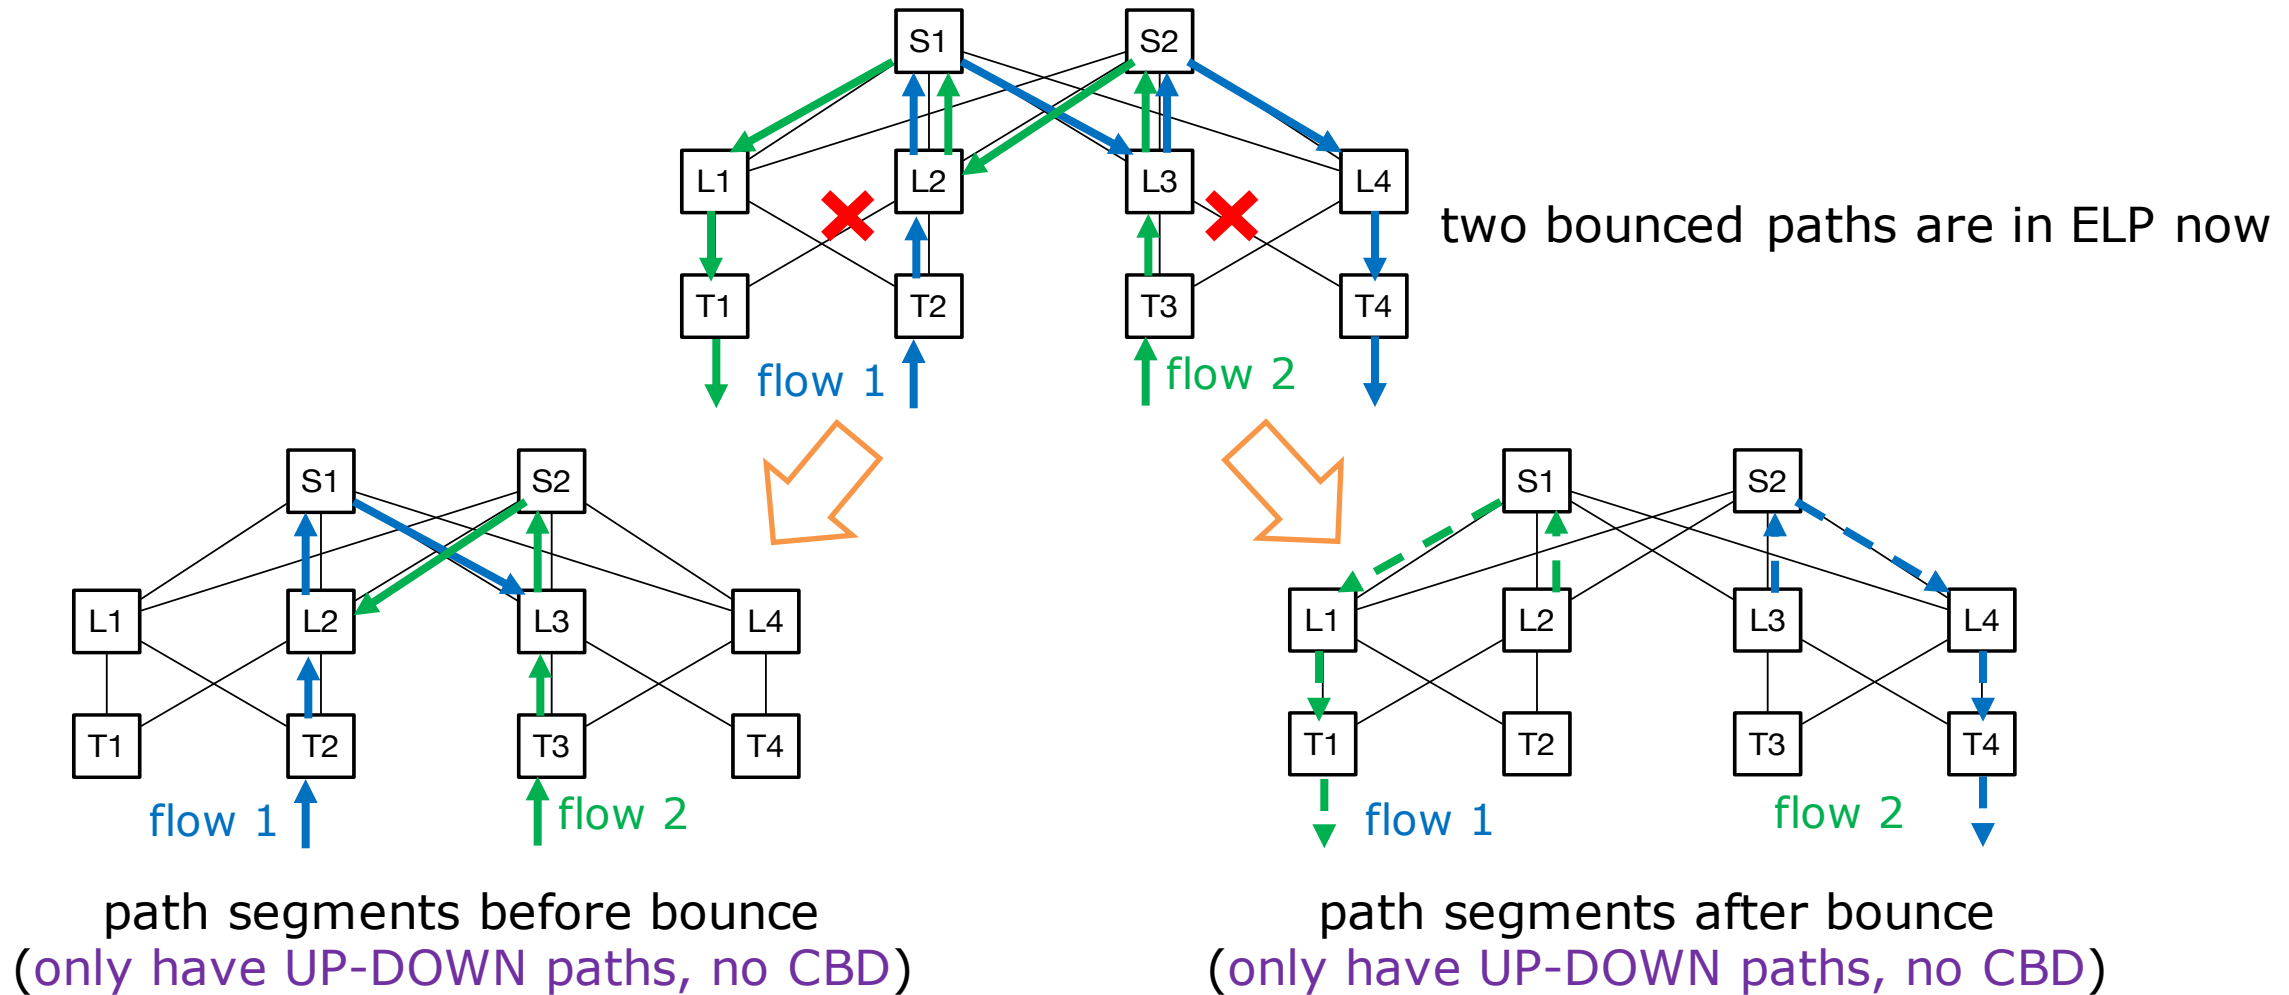
---




**ELP** = shortest paths + 1-bounce paths

(ELP has CBD now!)


# Segmenting ELP into CBD-free Subsets



# Isolating Path Segments with Tags





**tag 1** → path segments before bounce



**tag 2** → path segments after bounce




# Isolating Path Segments with Tags




Adding a rule at switch L3: (Tag = 1, Inport=S1, OutPort = S2) -> NewTag = 2


# No CBD after Segmentation



**tag 1**



**tag 2**




buffer dependency graph

packets with tag  $i \rightarrow i$ -th lossless queue

~~CBD: L2 > S1 > L3 > S2 > L2~~

# What If k-bounce Paths all in ELP?

---



**solution:** just segmenting ELP into k CBD-free subsets based on number of bounced times!

**ELP** = shortest up-down paths + ~~1-bounce paths~~  
**k**-bounce paths

# Summary: Tagger Design for Clos Topology

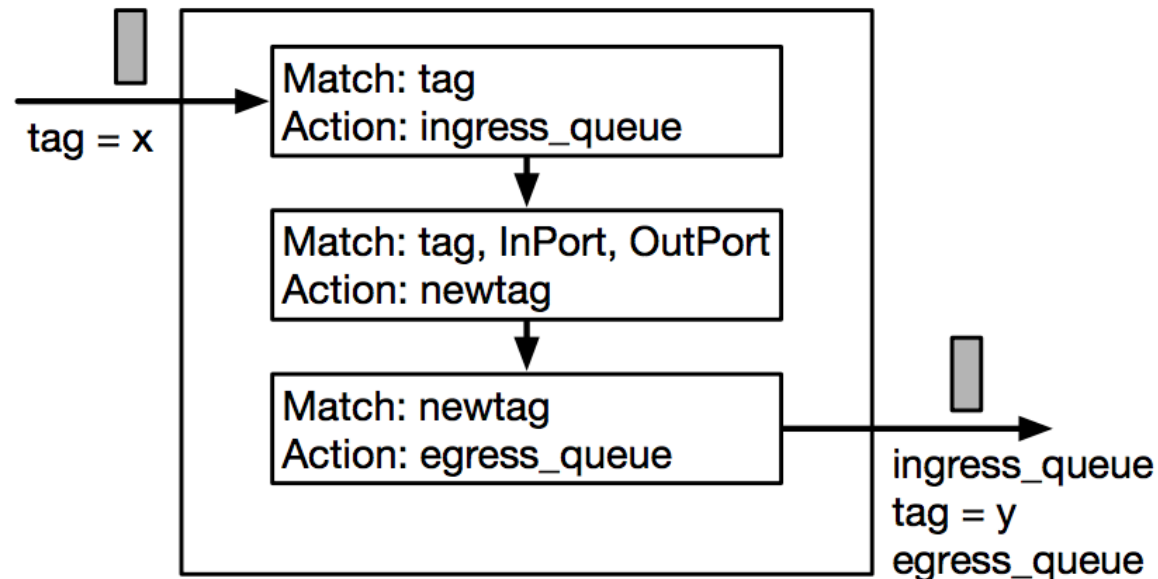
---

1. Initially, packets carry with tag = 1
2. pre-install match-action rules at switches:
  - DOWN-UP bounce: increase tag by 1
  - Enqueue packets with tag  $i$  to  $i$ -th lossless queue ( $i \leq k+1$ )
  - Enqueue packets with tag  $i$  to lossy queue ( $i > k+1$ )

For Clos topology, Tagger is optimal in terms of # of lossless priorities.

# How to Implement Tagger?

- DSCP field in the IP header as the tag carried in the packets
- build 3-step match-action pipeline with basic ACL rules available in commodity switches



# Tagger Meets All the Three Challenges

---


1. Work with existing routing protocols & hardware
2. Work with link failures & routing errors
3. Work with limited number of lossless queues

# More Details in the Paper


---

- Proof of Deadlock freedom
- Analysis & Discussions
  - Algorithm complexity
  - Optimality
  - Compression of match-action rules
  - ...


# Evaluation-1: Tagger prevents Deadlock



Scenario: two flows forms CBD



(a) Without Tagger



(b) With Tagger

Tagger avoids CBD caused by bounced flows, and prevents deadlock!



# Evaluation-2: Scalability of Tagger

---


Switches	Ports	Longest ELP	Lossless Priorities	Max Rules
100	32	5	2	40
500	64	6	3	76
1,000	64	6	3	88
2,000	64	7	3	98
2,000 (*)	64	7	4	135

\* last entry includes additional 20,000 random paths.


Match-action rules and priorities required for Jellyfish topology

Tagger is scalable in terms of number of lossless priorities and ACL rules.

# Evaluation-3: Overhead of Tagger



(a) Throughput



(b) Latency

Tagger rules have no impact on throughput and latency

# Conclusion

---

- Tagger: a tagging system guarantees deadlock-freedom
  - **Practical:**
    - require no change to existing routing protocols
    - implementable with existing commodity switching ASICs
    - work with limited number of lossless priorities
  - **General:**
    - work with any topologies
    - work with any ELPs

**Thanks!**