


Using RDMA Efficiently for Key-Value Services

Anuj Kalia (CMU)

Michael Kaminsky (Intel Labs), David Andersen (CMU)

RDMA



Remote Direct Memory Access:

A network feature that allows direct access to the memory of a remote computer.

HERD

1. Improved understanding of RDMA through micro-benchmarking
2. High-performance key-value system:
 - Throughput: 26 Mops (*2X higher than others*)
 - Latency: 5 μ s (*2X lower than others*)

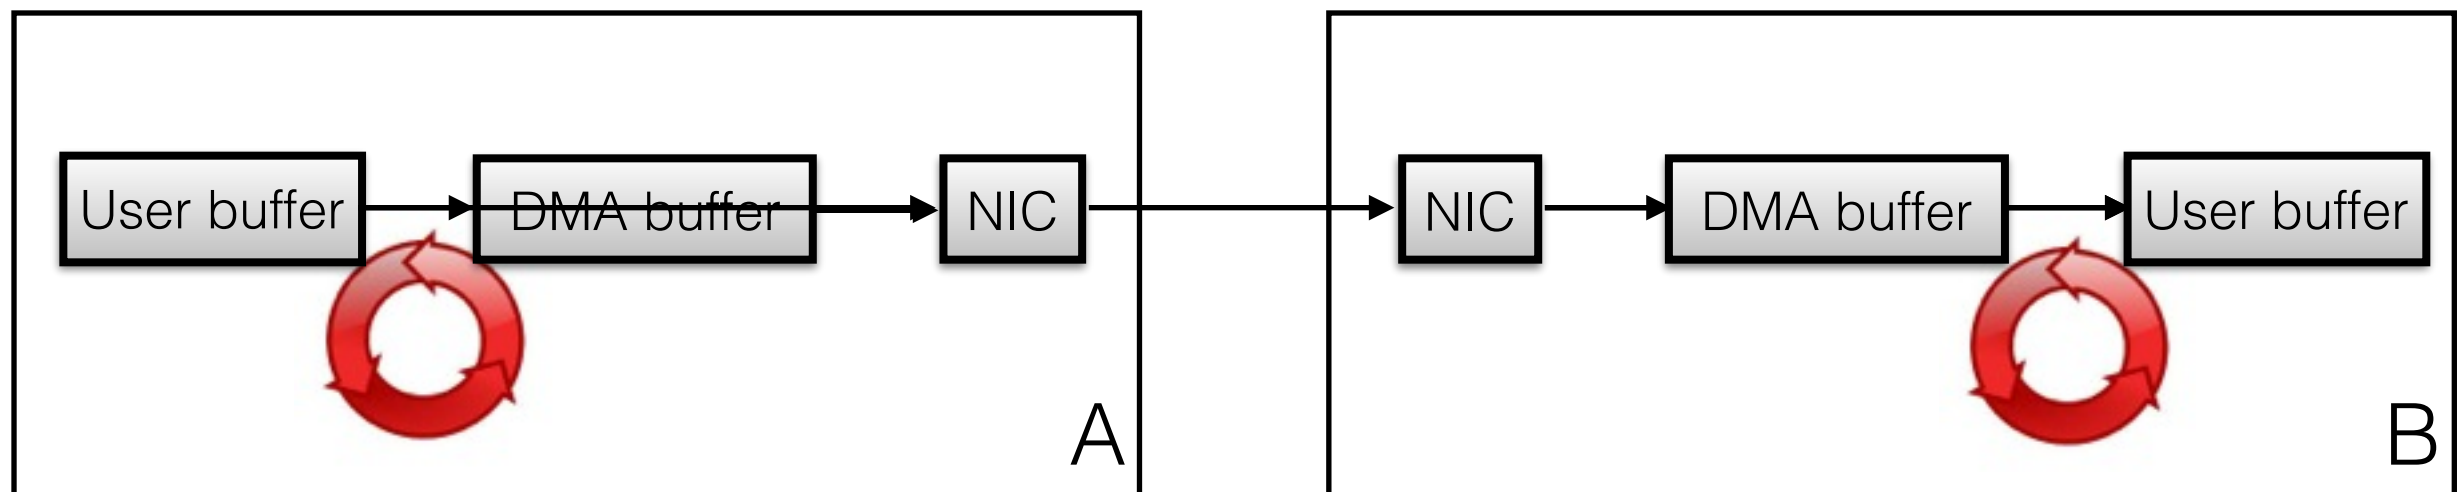
RDMA intro

Features:

- Ultra-low latency: 1 μ s RTT
- Zero copy + CPU bypass

Providers:


InfiniBand, RoCE,...




RDMA in the datacenter

48 port 10 GbE switches

Switch	RDMA	Cost
Mellanox SX1012	YES	\$5,900
Cisco 5548UP	NO	\$8,180
Juniper EX5440	NO	\$7,480



In-memory KV stores



Interface: GET, PUT

Requirements:

- Low latency
- High request rate

RDMA basics

Verbs


RDMA read:

```
READ(local_buf, size, remote_addr)
```

RDMA write:

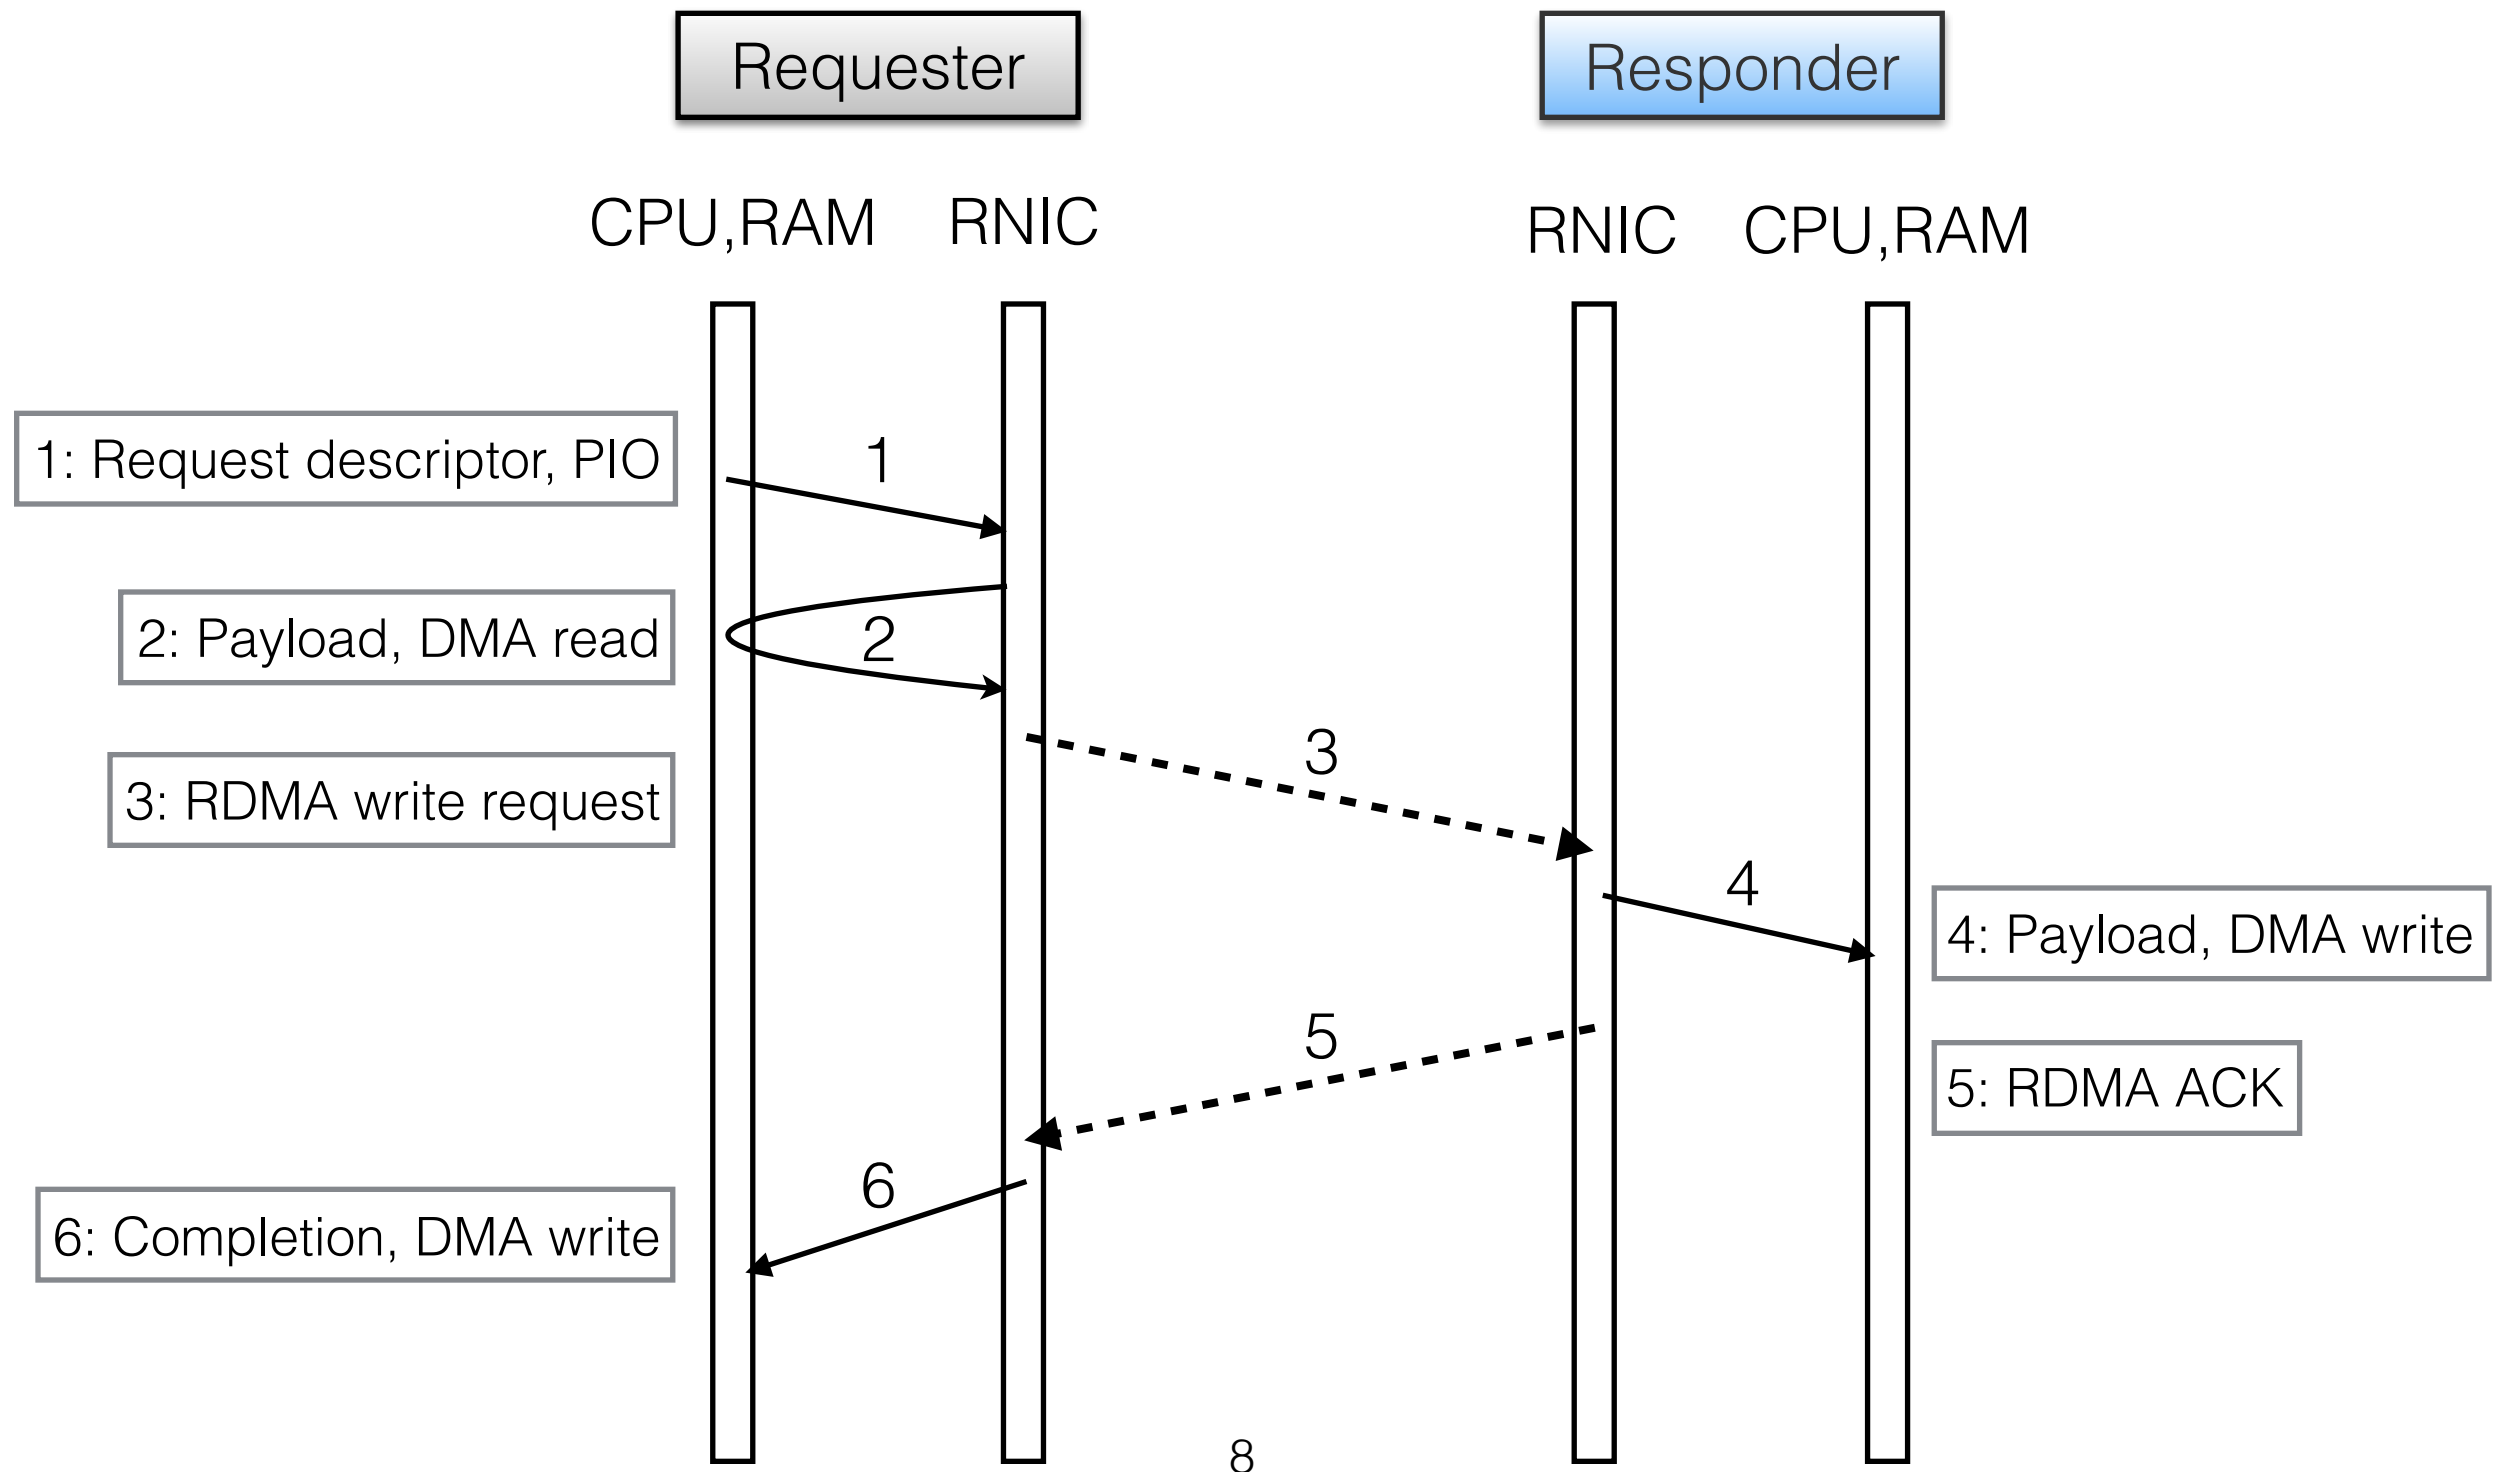
```
WRITE(local_buf, size, remote_addr)
```

⋮



RNIC

Life of a WRITE



Recent systems

Pilaf [ATC 2013]

FaRM-KV [NSDI 2014]: an example usage of FaRM


Approach: RDMA reads to access remote data structures

Reason: the allure of CPU bypass

The price of CPU bypass

Key-Value stores have an inherent level of indirection.

An index maps a keys to address. Values are stored separately.




At least 2 RDMA reads required:
 ≥ 1 to fetch address
1 to fetch value

Not true if value is in index


The price of CPU bypass



The price of CPU bypass




The price of CPU bypass



Our approach

Goal	Main ideas
#1: Use a single round trip	Request-reply with server CPU involvement + WRITES faster than READs
#2. Increase throughput	Low level verbs optimizations
#3. Improve scalability	Use datagram transport

#1: Use a single round trip



#1: Use a single round trip

Operation	Round Trips	Operations at server's RNIC
READ-based GET	2+	2+ RDMA reads
HERD GET	1	2 RDMA writes




Lower latency




High throughput

RDMA WRITES faster than READs



Setup: Apt Cluster

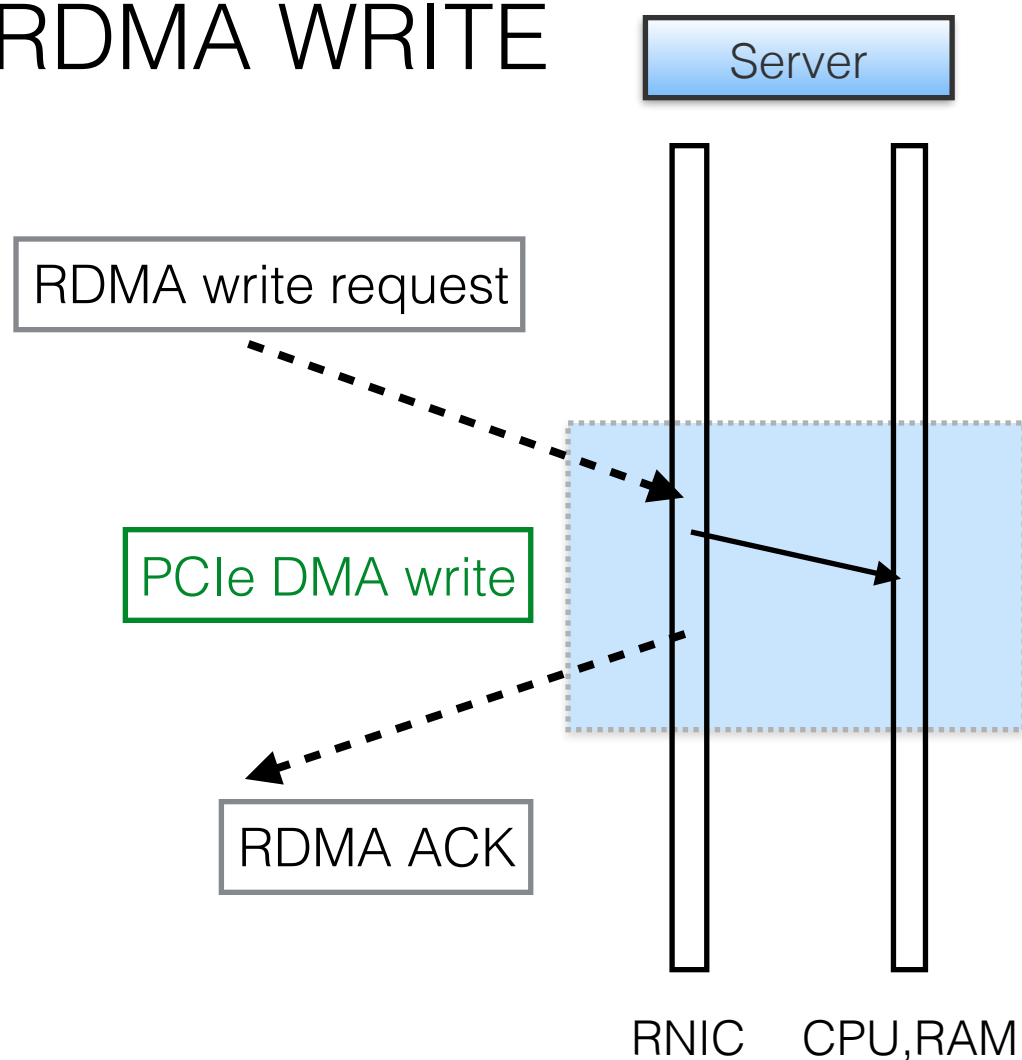
192 nodes, 56 Gbps IB



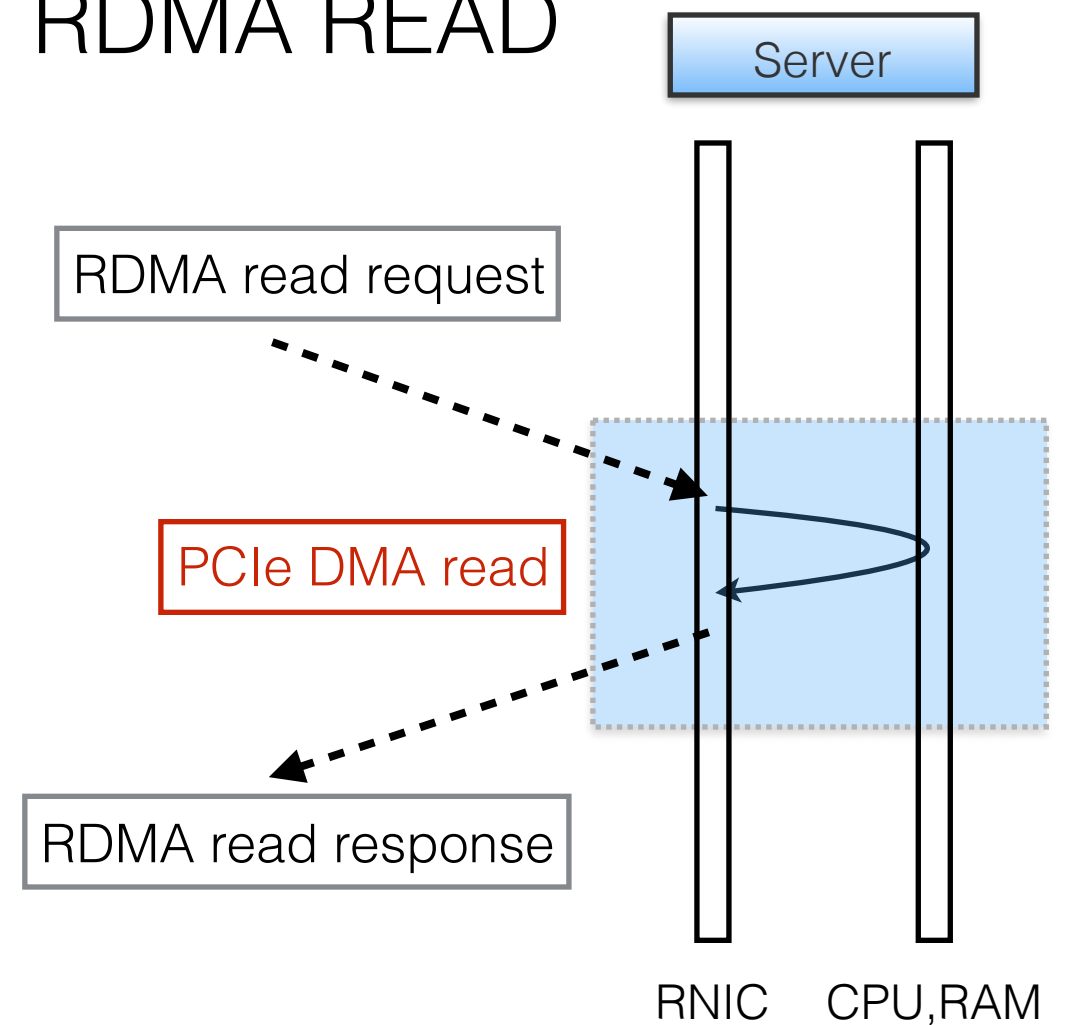
RDMA WRITES faster than READs

Reason: PCIe writes faster than PCIe reads

RDMA WRITE




RDMA READ




High-speed request-reply

Request-reply throughput:

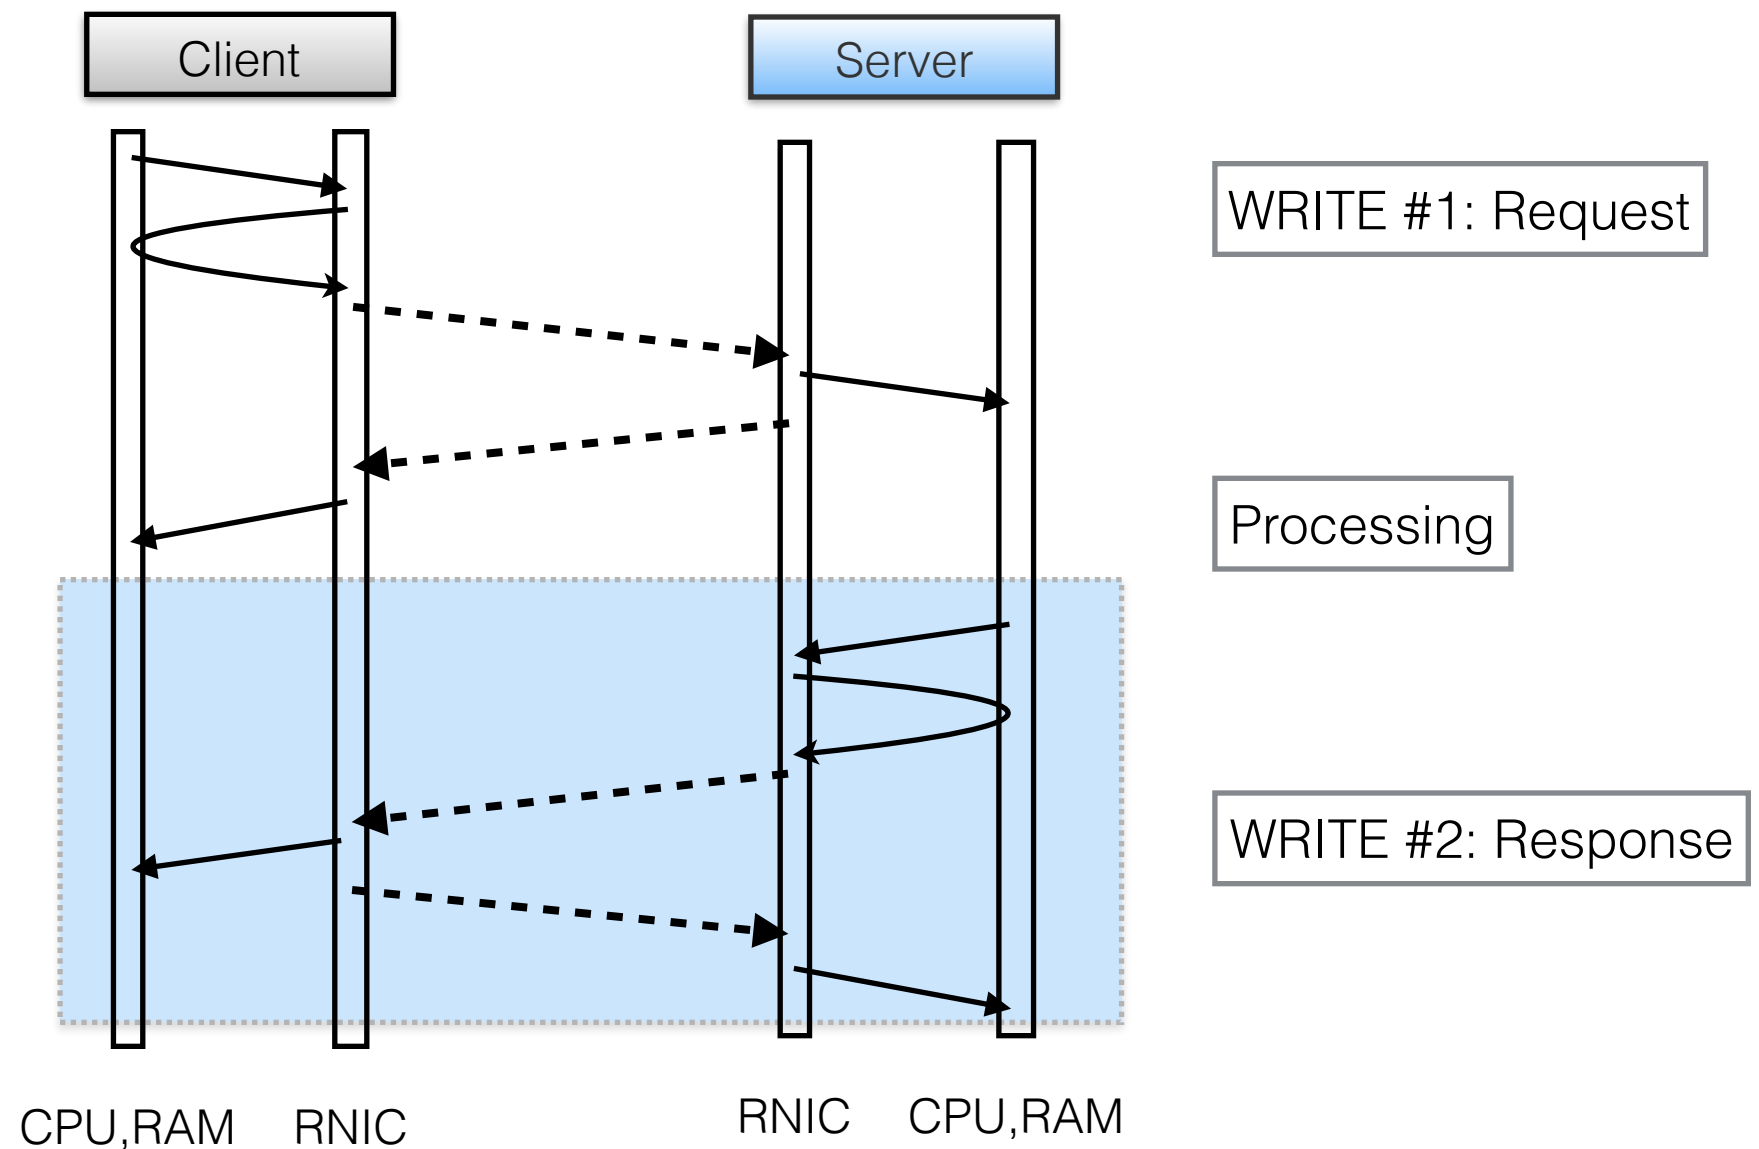


Setup: one-to-one client-server communication



#2: Increase throughput

Simple request-reply:

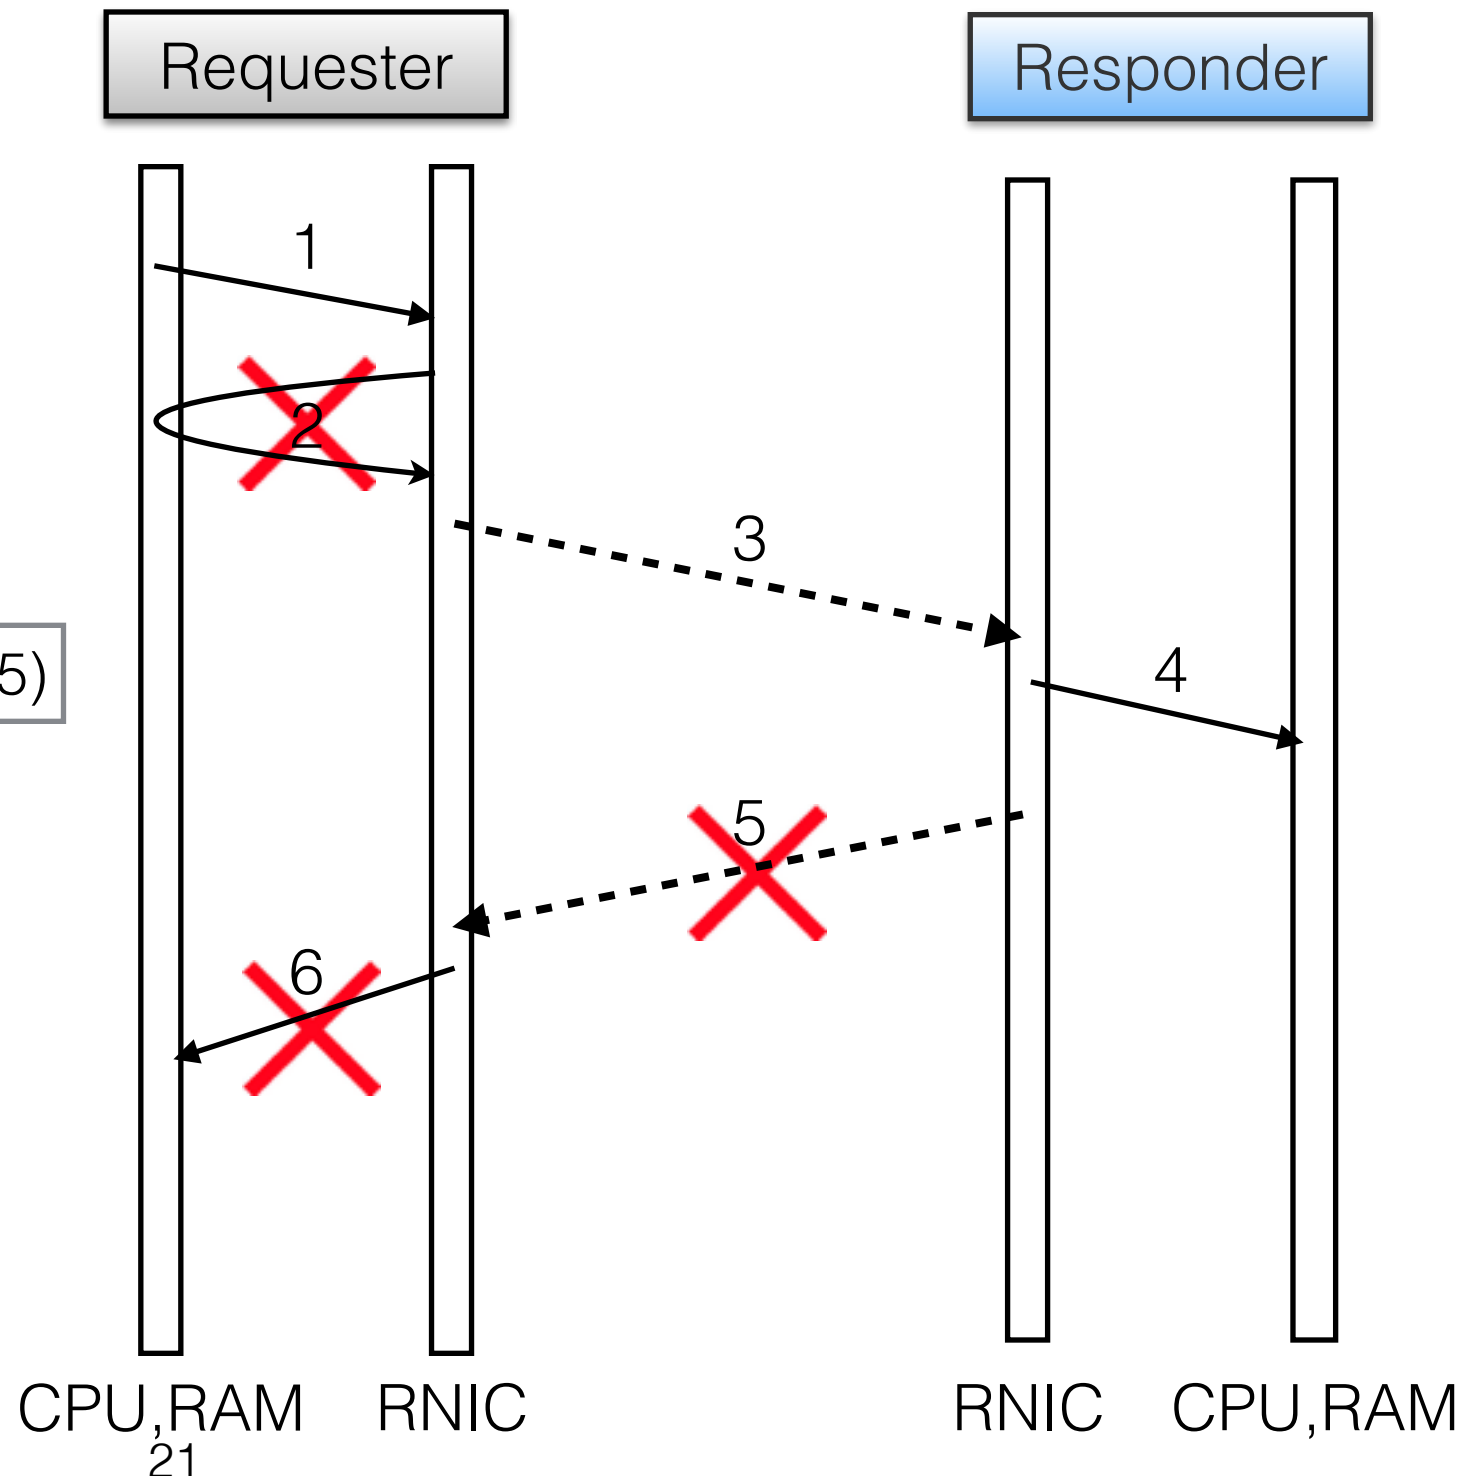


Optimize WRITES

+inlining: encapsulate payload in request descriptor (2→1)


+unreliable: use unreliable transport (- 5)

+unsignaled: don't ask for request completions (- 6)




#2: Increase throughput

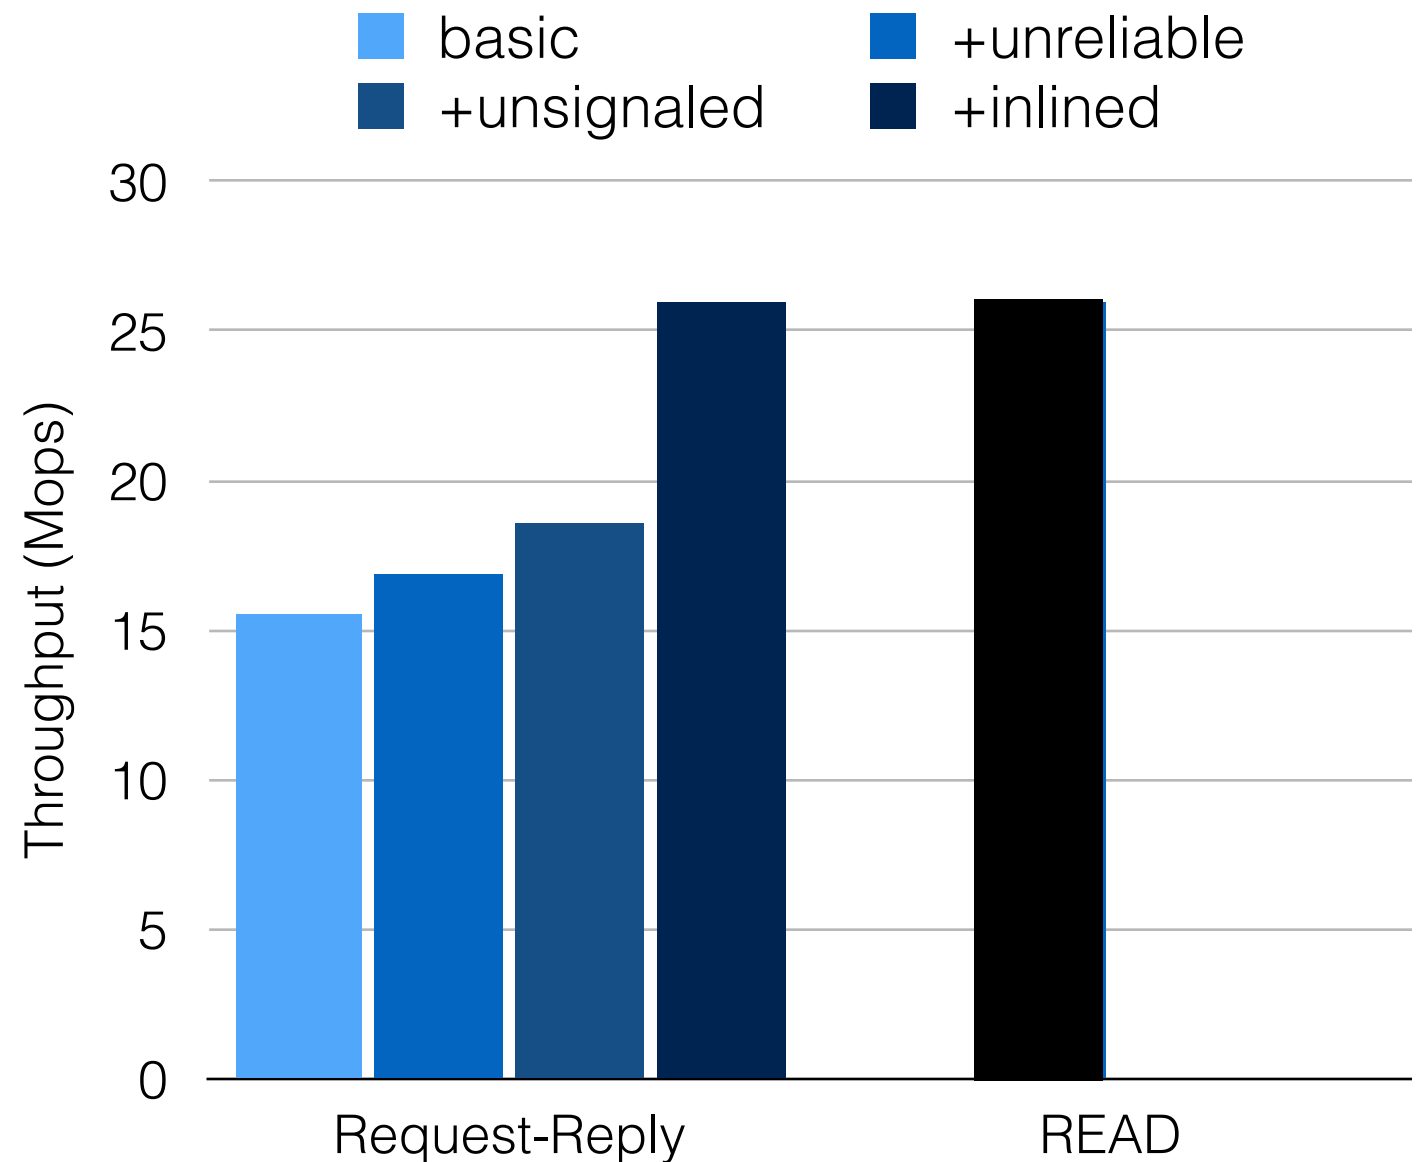
Optimized request-reply:




#2: Increase throughput




Setup: one-to-one client-server communication





#3: Improve scalability



Setup

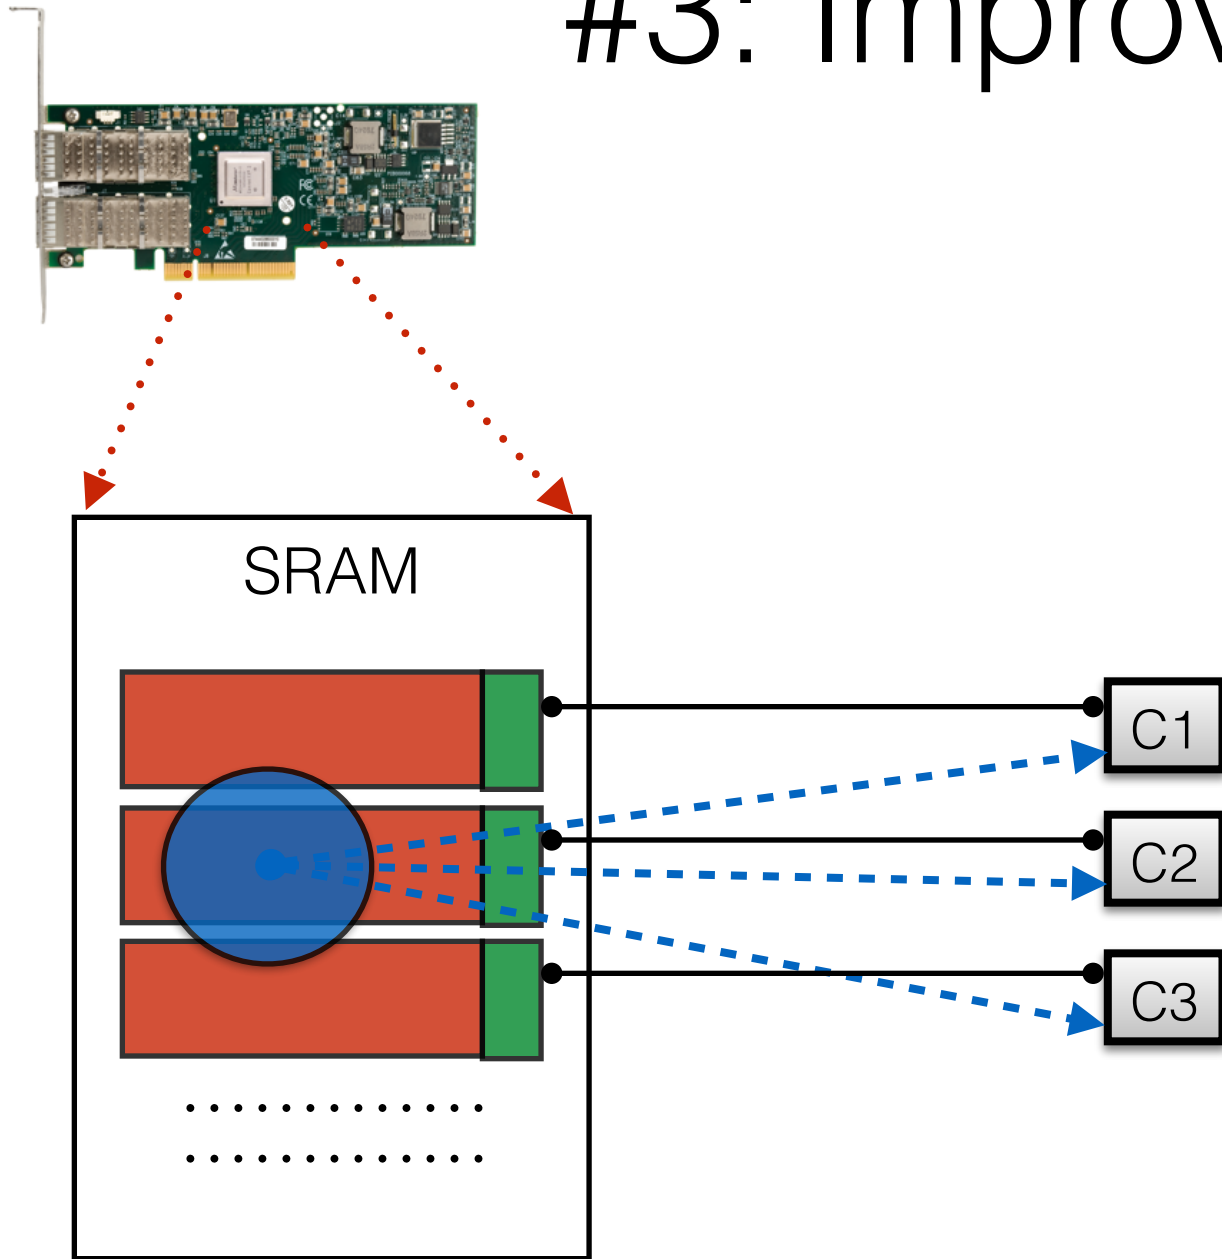




#3: Improve scalability



$||\text{state}|| > \text{SRAM}$

#3: Improve scalability



Inbound scalability \gg outbound *because*
inbound state () \ll outbound ()



Use datagram for outbound replies



Datagram only supports SEND/RECV.
SEND/RECV is slow.




SEND/RECV is slow only at the receiver




Scalable request-reply

→ RDMA write, connected
---→ SEND, datagram



Setup



Evaluation


HERD = Request-Reply + MICA [NSDI 2014]

Compare against emulated versions of Pilaf and FaRM-KV

- No datastore
- Focus on maximum performance achievable


Latency vs throughput

48 byte items, GET intensive workload




Latency vs throughput

48 byte items, GET intensive workload



Throughput comparison

16 byte keys, 95% GET workload




HERD

- Re-designing RDMA-based KV stores to use a single round trip
 - WRITES outperform READs
 - Reduce PCIe and InfiniBand transactions
 - Embrace SEND/RECV
- Code is online: <https://github.com/efficient/HERD>

Throughput comparison

16 byte keys, 95% GET workload



Throughput comparison

